Zaheer, M., Guruganesh, G., Dubey, K. A., Ainslie, J., Alberti, C., Ontanon, S., Pham, P., Ravula, A., Wang, Q., Yang, L., et al. Big bird: Transformers for longer sequences. *Advances in Neural Information Processing Systems*, 33:17283–17297, 2020.

Zhang, Z., Sheng, Y., Zhou, T., Chen, T., Zheng, L., Cai, R., Song, Z., Tian, Y., Ré, C., Barrett, C., et al. H_2O: Heavy-hitter oracle for efficient generative inference of large language models. *arXiv preprint arXiv:2306.14048*, 2023.

A. Detailed Results

Figures A1 to A3 report the compression/performance trade-off curves for all models and tasks that were evaluated.

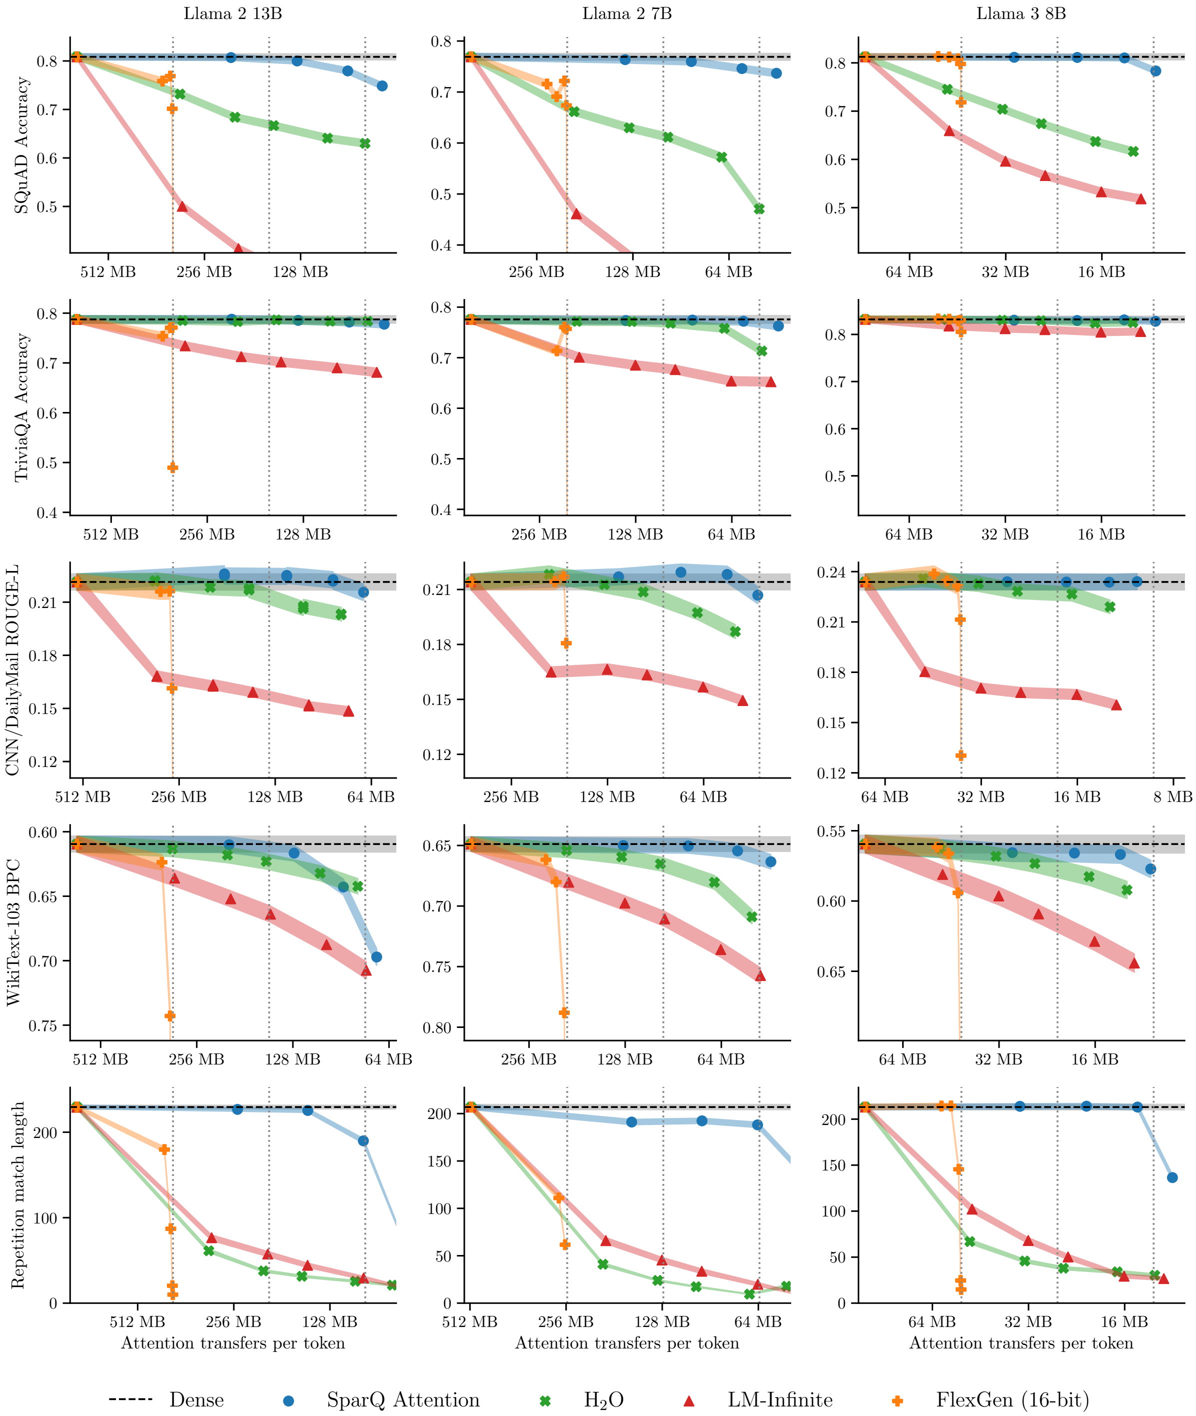

Figure A1: Compression versus performance trade-off curves over all tasks for the Llama 2 and Llama 3 model families. The y-axis minimum is set to $(0.5, 0.5, 0.5, 1.25, 0.0)\times$ the dense baseline for the tasks, reading top-to-bottom, in order to give a consistent view of the performance loss across models. Vertical dotted lines show $1/2\times$, $1/4\times$ and $1/8\times$ compression versus dense. Shaded lines show ±1 standard error of the mean (uncertainty due to a finite test set).

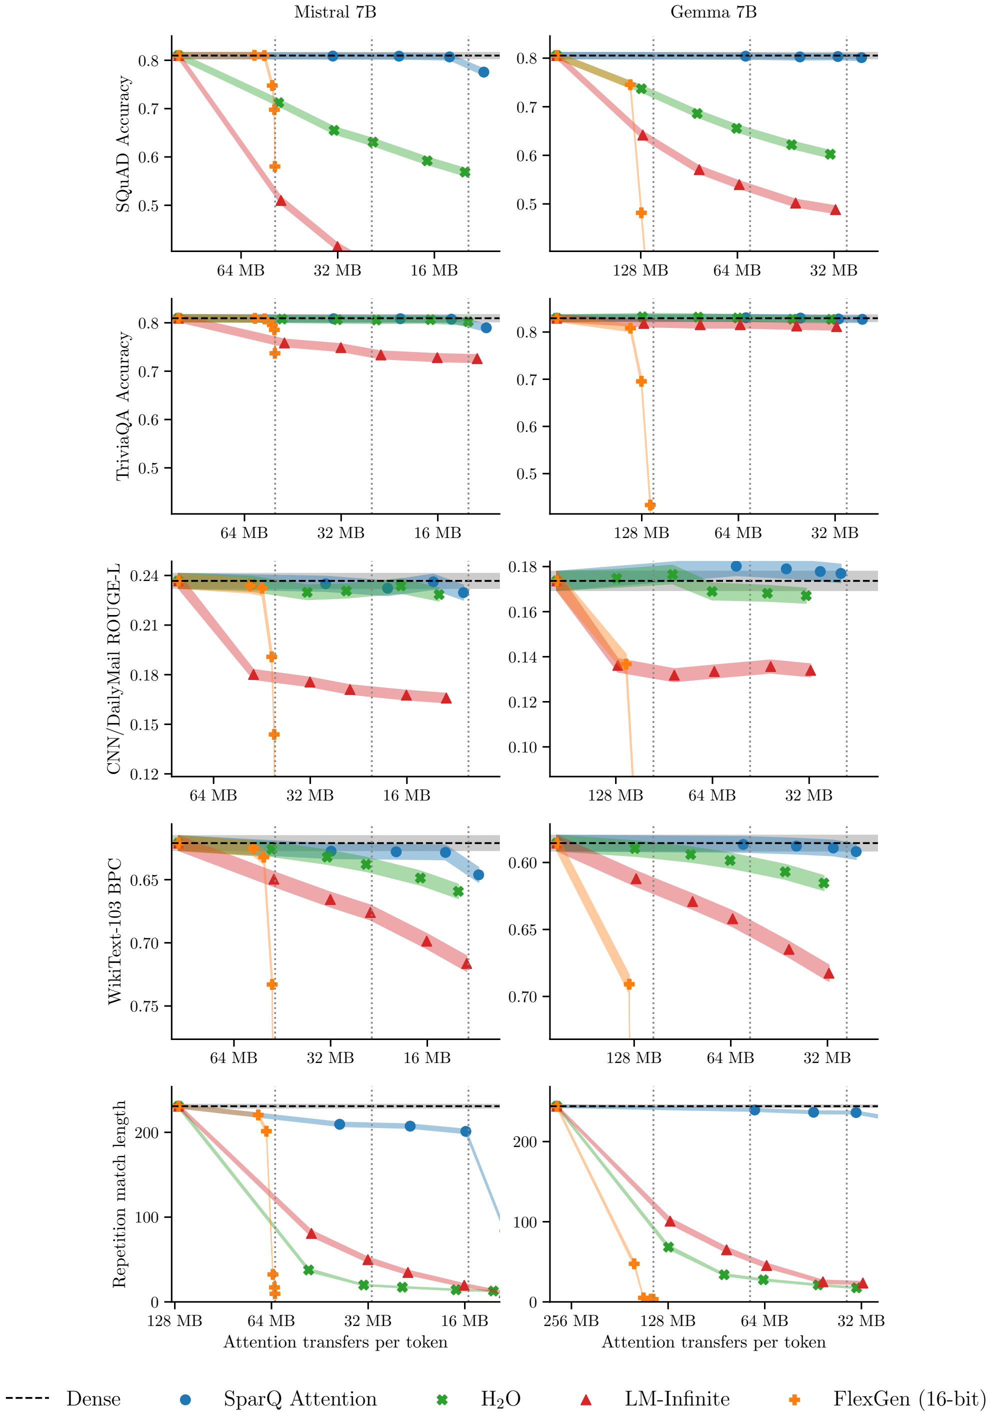

Figure A2: Compression versus performance trade-off curves over all tasks for Mistral 7B and Gemma 7B models. The y-axis minimum is set to $(0.5, 0.5, 0.5, 1.25, 0.0)\times$ the dense baseline for the tasks, reading top-to-bottom, in order to give a consistent view of the performance loss across models. Vertical dotted lines show $1/2\times$, $1/4\times$ and $1/8\times$ compression versus dense. Shaded lines show ± 1 standard error of the mean (uncertainty due to a finite test set).